

RESEARCH

The AI Market and the Road Ahead for an Experienced Engineer: A 12-24 Month Action Map

Contents

01	Executive Summary	4
02	Background	7
03	Part A — Market State and Where the Puck Is Going [HIGH confidence]	10
04	Part B — How Software Engineers Are Actually Evolving [HIGH confidence]	13
05	Part C — The 12–24 Month Career Roadmap for a Senior Generalist [HIGH confidence on direction; MEDIUM on sequencing]	17
06	Part D — Where Business Gets Value, and Where a Very Small Company Can Leverage AI [HIGH confidence]	21
07	Recommendations	25
08	Limitations	28
09	Sources	31

Research ID: R-00144

Date: 2026-06-23

Mode: deep

Depth: deep

Primary Question: Given the current AI market trajectory, where should a 44-year-old engineer with 20 years of experience place their bets over the next 12–24 months — and where can a very small company realistically extract AI value today?

IW Innovation Ways

CHAPTER 01

Executive Summary

CHAPTER 01

Executive Summary

The AI market in mid-2026 is enormous, concentrated, and — at the infrastructure layer — financially stretched, but the *application* layer where an experienced engineer and a small company actually operate is healthier and clearer than the headlines suggest. Three signals dominate. First, the value is real but narrow: coding assistance leads enterprise adoption by ~10x over every other use case, followed by customer support and internal search ([a16z](#)), while ~95% of broad GenAI pilots still show no P&L impact ([MasterOfCode](#)) and fewer than 40% of organizations report EBIT impact ([McKinsey](#)). The differentiator between the winners and the 95% is **workflow redesign**, not tool count. Second, the job market is bifurcating sharply by seniority: entry-level/programmer roles are contracting (programmer employment -27.5% 2023–2025; new grads down to ~7% of Big Tech hires from 32% in 2019) while engineers 30+ in high-AI-exposure roles *grew* 6–12% ([IEEE Spectrum](#); [finalroundai](#)). Experience is now an asset, not a liability — judgment is the scarce input. Third, AI is an **amplifier of existing engineering maturity**, not an autonomous replacement: DORA 2025 found 90% AI adoption finally correlating with higher delivery throughput, but only for teams with strong platforms and workflows ([DORA 2025](#)).

For a senior generalist, the 12–24 month bet is **not** to become an ML researcher. It is to become the person who designs, evaluates, and operates reliable AI-powered systems — RAG, agents/tool-calling (MCP), evaluation/observability, cost/latency, and guardrails — exactly the stack our prior research (R-00130/R-00134) found is screened for across thousands of postings, and exactly the judgment-heavy work that 20 years of production experience uniquely qualifies you for. For a very small company, the highest-ROI move is twofold: (1) use AI internally to compress your own delivery and ops (proven, low-risk, fast payback — SMBs report saving 20+ hours and €500–€2,000/month), and (2) productize a **vertical micro-SaaS or AI-enabled service** where domain depth — not model access — is the moat. The macro risk (a capex bubble; ~\$660–690B hyperscaler spend in 2026 against thin revenue) argues for staying at the application/services layer and avoiding capital-intensive bets.

IW Innovation Ways

CHAPTER 02

Background

CHAPTER 02

Background

This is the career-and-SMB-leverage installment of an AI research series. It explicitly **builds on and cross-references** seven priors filed 2026-06-13:

- **R-00129** — AI market 2026 landscape (provider share, pricing collapse, agentic shift, bubble risk). Treated here as the macro baseline; this report confirms its trajectory with mid-2026 data and extracts the *implications* for an individual and a small firm.
- **R-00130** — SWE → app-layer AI engineer pivot (the screened-for skill stack). The canonical skills source; this report layers a *sequenced roadmap* and a *de-prioritization list* on top.
- **R-00131 / R-00135** — Innovation Ways GTM and EU services pricing/packaging. Referenced for the "sell AI services" path; this report adds the *internal-leverage* and *productization* paths those did not cover.
- **R-00132** — trending solo AI showcase projects (portfolio signal).
- **R-00133** — agent eval/observability tooling (the technical depth behind "evaluation is the shipping gate").
- **R-00134** — EU AI-engineer resume teardown (how to present the pivot).

Gaps this research fills: (a) an integrated, sequenced 12–24-month career roadmap with an explicit *what-to-de-prioritize* list; (b) where businesses get *proven* value and where a *very small company* can leverage AI **internally** (not just sell it); (c) a mid-2026 reality check on the market and the bubble risk as it bears on personal/SMB strategy.

CHAPTER 03

Part A — Market State and Where the Puck Is Going [HIGH confidence]

CHAPTER 03

Part A — Market State and Where the Puck Is Going [HIGH confidence]

The macro picture from R-00129 holds, and the concentration has intensified. Hyperscaler capex for 2026 is now tracking ~\$660–690B combined (Amazon ~\$200B, Alphabet \$175–185B, Meta \$115–135B, Microsoft ~\$120B, Oracle ~\$50B) ([Futurum](#); [tooldirectory](#)). The revenue under that spend remains thin and circular: one widely cited figure has OpenAI spending ~\$60B/yr on compute against ~\$13B revenue ([tooldirectory](#)), and the US Federal Reserve flagged AI as a top systemic risk in 2026 ([Man Group](#)). Man Group's structural warning from R-00129 still bites: token costs falling ~70%/year would need ~225%/year demand growth just to hold economics flat ([Man Group](#)).

Strategic read for an individual/SMB: the bubble risk is concentrated in the *infrastructure and frontier-model* layer — GPUs, data centers, model training. The *application and services* layer benefits from the same dynamics that threaten the infra layer: collapsing token prices (>90% since 2023, per R-00129) make AI-powered products cheaper to run every quarter. **Build on top, not underneath.** A small firm should treat models as a deflationary, swappable commodity (R-00129: only ~11% of enterprises switch providers annually, but switching is technically cheap) and avoid any bet whose economics depend on cheap capital or proprietary compute.

The agentic shift is now genuinely in production, but unevenly. 31% of enterprises have at least one agent in production (banking/insurance leading at 47%, healthcare/government trailing at 18%), and ~80% of enterprise apps shipped/updated in Q1 2026 embed at least one agent, up from 33% in 2024 ([Digital Applied](#)). Median time-to-value on agent deployments is ~5.1 months, but Gartner-style caution persists: 40%+ of agentic projects are at risk of cancellation by 2027, with integration into existing systems (46%) the top blocker ([Digital Applied](#)). **MCP has consolidated as the integration standard** (9,400+ public servers), confirming R-00130's "USB-C for AI" framing. The takeaway: agents are real, but the hard, valuable, defensible work is *integration, reliability, and evaluation* — engineering, not prompting.

CHAPTER 04

Part B — How Software Engineers Are Actually Evolving [HIGH confidence]

CHAPTER 04

Part B — How Software Engineers Are Actually Evolving [HIGH confidence]

The job market is bifurcating by seniority, and experience is winning. This is the single most important — and most encouraging — finding for a 44-year-old engineer:

- **Junior/programmer roles are contracting.** Programmer employment fell 27.5% in the US between 2023 and 2025; entry-level hiring at 15 major tech firms fell 25% (2023 → 2024); new grads are ~7% of Big Tech hires vs 32% in 2019; and 22–25-year-old developer employment fell ~20% from its late-2022 peak ([IEEE Spectrum](#); [finalroundai](#)).
- **Senior/experienced roles are growing.** For workers 30+ in the highest AI-exposure categories, employment grew 6–12% from late 2022 to May 2025 ([finalroundai](#)). Notably, the more *design-oriented* "software developer" title fell only 0.3% while "programmer" fell double digits ([IEEE Spectrum](#)) — the market is repricing *coding-as-typing* down and *judgment/design* up.
- **The pyramid is compressing, not collapsing.** AI skills now appear in 42% of software job descriptions (up from 8% in 2022); AI/ML-engineer postings are up ~85% YoY with 20–30% salary premiums; and the BLS still projects 17% software-employment growth through 2033 ([finalroundai](#)). The "demise" narrative is overstated ([CNN](#)).

The role itself is shifting from execution to judgment. Across sources the consistent message is that the senior engineer's value "is no longer measured by how much code they personally write," but by designing the environment where AI-assisted development happens safely — deciding *what* to build, *how it behaves under pressure*, and *where AI cannot be trusted* ([Security Boulevard](#)). Addy Osmani frames the next two years as engineers becoming orchestrators and reviewers of AI output, with verification skill as the bottleneck ([Addy Osmani](#)).

AI is an amplifier of engineering maturity, not a substitute for it. DORA 2025 is the most rigorous data point: AI adoption among developers hit 90% (median ~2 hours/day), 80%+ report productivity gains, and — reversing 2024's finding — AI now correlates with *higher* delivery throughput. But the report's central thesis is the "amplifier effect": AI multiplies the conditions it finds, "strengthening high-performing teams while exposing weaknesses" in those with fragmented processes; a high-quality internal platform is the strongest predictor of unlocking AI value ([DORA 2025](#)). For a 20-year veteran, this is the core insight: **your accumulated systems/platform/quality judgment is precisely the multiplier AI needs.**

The tooling baseline has standardized. The market settled into a *composable multi-tool stack*, not a single winner: developers run ~2.3 AI coding tools on average, commonly Claude Code + Cursor for different parts of the workflow, with OpenCode reaching ~7.5M developers ([The New Stack](#); [Developers Digest](#)). Fluency with agentic coding tools is now table-stakes, not a differentiator.

Emerging specializations (where senior engineers are concentrating):

1. **AI/LLM application engineering** — RAG, agents, tool-calling, the R-00130 stack. Largest demand pool.
2. **AI evaluation & observability ("evals")** — the shipping gate; R-00133's domain. Under-

supplied relative to demand.

3. **Platform / AI infrastructure engineering** — internal platforms, orchestration, the DORA "multiplier."

4. **Security engineering for AI** — prompt-injection defense, guardrails, trust boundaries; information-security roles grew double digits while programming fell ([IEEE Spectrum](#)).

5. **Domain-embedded AI** (legal, health, fintech) — where vertical depth + AI compounds.

CHAPTER 05

**Part C — The 12–24 Month
Career Roadmap for a Senior
Generalist [HIGH confidence on
direction; MEDIUM on
sequencing]**

CHAPTER 05

Part C — The 12-24 Month Career Roadmap for a Senior Generalist [HIGH confidence on direction; MEDIUM on sequencing]

The strategic frame: **do not chase model-building; become the reliability-and-systems owner of AI products.** R-00130 established that app-layer AI engineering "rewards strong software engineers more than ML researchers" — math depth is largely a myth for this track; production instincts are the differentiator. Your 20 years *is* the qualification. Below is a sequenced, build-in-public plan.

What to INVEST in (ranked by leverage)

Priority	Skill area	Why it pays off	Source anchor
1	RAG done properly (chunking, embeddings, hybrid search, re-ranking; pgvector → Pinecone at scale)	Single most in-demand skill (~36–39% of postings); depth separates you from tutorial-completers	R-00130; R-00134
2	Evaluation & observability (offline + online evals, LLM-as-judge, tracing, regression gates)	"If you can't measure your agent, you can't ship it" — the shipping gate, under-supplied	R-00130; R-00133
3	Agents & tool-calling + MCP (state machines, checkpointing, tool design)	Fastest-growing category; "bad tool design causes 70% of agent failures"	R-00130; Digital Applied
4	Context engineering (managing everything the model sees, not just prompts)	The 2026 framing that superseded prompt engineering	R-00130
5	Cost/latency optimization (caching, routing, batch APIs, streaming)	"Most underrated skill"; cuts bills 40–70%; commands premiums	R-00130
6	AI security & guardrails (OWASP LLM risks, prompt-injection, least privilege)	A growing, defensible specialization; treated as classic security engineering	R-00130; IEEE Spectrum
7	Agentic-coding fluency (Claude Code/Cursor/OpenCode as a daily multiplier)	Table-stakes; also makes you 10–20x more productive as a solo builder	The New Stack; a16z

What to DE-PRIORITIZE (deliberately, to free time) [MEDIUM confidence]

- **Don't try to become an ML researcher / model trainer.** <2% of AI-engineer postings touch classical ML/CV; fine-tuning appears in only ~8.5% (R-00130). Deep math/training is the wrong track for your background.
- **Don't over-index on any single framework.** No framework dominates (LangChain ~18.8%); "architectural understanding matters most" (R-00130). Learn one, but treat them as swappable.
- **Don't invest in skills AI is commoditizing:** boilerplate/CRUD coding, syntax memorization, manual-only QA execution, style-only code review, boilerplate doc generation ([gnxtsystems](#)). These are exactly what AI now does well; pivot the time saved into design/review/evals.
- **Don't collect certifications or tutorial notebooks as your portfolio** — they read as red flags; a *shipped, measured* LLM feature is the signal (R-00130/R-00134).
- **Don't bet on infra/frontier-layer plays** (training compute, foundation models) — capital-intensive and bubble-exposed (Part A).

The roadmap, sequenced

- **Months 0–3 — Re-title and ship one thing.** Reframe existing production/scaling/API experience as "LLM systems in production" (R-00134's "re-titling, not buzzword pivot"). Build **one** end-to-end RAG-or-agent feature on a real problem (ideally inside your own company), instrumented with evals, latency, and cost numbers. This single artifact is worth more than any course.

- **Months 3–9 — Add evals + agents depth.** Layer evaluation/observability (R-00133 tooling) and a genuine tool-calling/MCP agent with state and guardrails onto that project. Write up the *measured* results publicly (the R-00132 "build in public" signal). Adopt agentic coding tools as your daily driver to compress your own throughput.
- **Months 9–18 — Specialize + monetize.** Pick a vertical you already understand and go deep (domain-embedded AI is where your 20 years compounds). For Innovation Ways this doubles as GTM (R-00131/R-00135): turn the capability into a productized assessment/POC offer.
- **Months 18–24 — Operate at the platform/architecture altitude.** Position as the person who designs the *AI-native development environment* and reliability gates for a team — the highest-value, lowest-substitution role (Security Boulevard; DORA "multiplier"). This is the defensible end-state for a senior generalist.

Why this fits a 44-year-old with 20 years specifically: the market is paying a premium for exactly what you have (judgment, systems thinking, production scars, reliability instinct) and discounting exactly what juniors have (raw coding throughput, now commoditized). You are not swimming against the current — the current is moving toward you. The risk is *not adapting the framing*, not your age or experience.

CHAPTER 06

Part D — Where Business Gets Value, and Where a Very Small Company Can Leverage AI [HIGH confidence]

CHAPTER 06

Part D — Where Business Gets Value, and Where a Very Small Company Can Leverage AI [HIGH confidence]

Enterprise reality is a barbell: ~95% of broad pilots show no P&L impact ([MasterOfCode](#)) and <40% report EBIT impact ([McKinsey](#)), yet a minority extract large, repeatable value. The dividing line is **workflow redesign**: McKinsey's high performers are ~3x more likely to have fundamentally redesigned workflows, the strongest predictor of EBIT impact ([McKinsey](#)). AI applied to isolated tasks fails; AI embedded in a redesigned process pays.

The proven-value map (where ROI is real)

The use cases with the clearest ROI (a16z's adoption barbell):

1. **Coding/developer productivity** — leads by ~10x; portfolio companies cite 10–20x productivity for top engineers ([a16z](#)).
2. **Customer support** — clear SOPs, high volume, transactional; measurable on tickets/CSAT/resolution ([a16z](#)).
3. **Internal search / knowledge discovery** — incl. vertical search (legal=Harvey at ~\$200M ARR in 3 years; health scribing) ([a16z](#)).
4. **Back-office cost takeout** — 26–31% reported cost savings in supply chain/procurement, finance, and people ops ([futransolutions/IBM-class reporting](#)) [MEDIUM].

Where a very small company can leverage AI today — two distinct paths

SMB adoption is now mainstream and *measurably productive*: ~58% of small businesses use GenAI (up from 40% in 2024); 80%+ of AI-using SMBs report productivity gains; many save 20+ hours/month and €/\$500–2,000/month; the average SMB uses ~5 AI tools and the SMB-vs-large adoption gap shrank from 1.8x to 1.2x ([AdAI](#); [Goldman Sachs](#); [business.com](#)). Small size is now an *advantage*: no legacy process to unwind, so you can redesign workflows AI-first from scratch — the exact lever McKinsey says enterprises struggle with.

Path 1 — Internal leverage (do this first; low-risk, fast payback). Use AI to compress your own delivery and operations:

- **Engineering throughput**: agentic coding (Claude Code/Cursor/OpenCode) as the daily stack — the single highest-ROI internal use, and you build the very skills Part C rewards. (This repo, IW AI Core, is itself an example of AI-orchestrated delivery.)
- **Go-to-market & ops**: content/marketing generation (the #1 SMB use case), proposal/spec drafting, customer-support assist, internal RAG over your own docs/contracts/code, finance/admin automation. SMBs report the biggest time savings exactly here.
- **Workflow redesign, not tool-bolting**: pick one process (e.g., proposal → POC → delivery) and rebuild it AI-native end-to-end. That is what converts pilots into P&L.

Path 2 — Productize (where the upside is, but choose carefully). The 2026 consensus is that the durable small-company product play is **vertical, AI-native micro-SaaS or AI-enabled services**, not horizontal "AI wrappers":

- The opportunity is "solve one specific problem for one specific industry, better than anyone, powered by AI"; a domain-expert solo/small founder can reach \$300–500K ARR in 12–18 months, building/launching in 4–12 weeks ([Growtoria](#); [AI Magicx](#)).
- **The moat is domain data/workflow/relationships, not model access** — "90% of AI wrapper startups will fail by end of 2026, 60–70% generate zero revenue" ([Growtoria](#)). Your 20 years of domain context is the defensibility.

- The opportunity is "solve one specific problem for one specific industry, better than anyone, powered by AI"; a domain-expert solo/small founder can reach \$300–500K ARR in 12–18 months, building/launching in 4–12 weeks ([Growtoria](#); [AI Magicx](#)).
- **The moat is domain data/workflow/relationships, not model access** — "90% of AI wrapper startups will fail by end of 2026, 60–70% generate zero revenue" ([Growtoria](#)). Your 20 years of domain context is the defensibility.
- The **services-firm path** (assessment → capped POC → milestone build → retainer) is fully specified in R-00131/R-00135, including EU benchmark rates and an EU-sovereignty premium of 20–40%. Path 2 and the services path can feed each other: deliver client POCs, harvest the repeatable one into a product.

The EU regulatory angle (relevant because you operate in the EU) [HIGH confidence]

The EU AI Act is now partly in force: prohibited practices since Feb 2025, GPAI-model obligations since Aug 2025; the **2 Aug 2026** deadline brings transparency obligations (applicable to all) and Annex III high-risk rules — though the *Digital Omnibus* (provisional agreement 7 May 2026, pending adoption) defers high-risk Annex III to **2 Dec 2027** ([EU Commission](#); [artificialintelligenceact.eu](#)). Crucially for a small firm: the Act deliberately *favours SMEs* — priority free access to regulatory sandboxes, simplified technical-documentation forms, proportionate (lower-bound) fines, and fees scaled to firm size ([artificialintelligenceact.eu](#)). 60%+ of EU SMEs have not started compliance ([Delbion](#)) — which is itself a **services opportunity** (AI-Act readiness/compliance for SMEs) that compounds with the R-00135 sovereignty premium.

Recommendations

CHAPTER 07

Recommendations

For you (the engineer), next 12–24 months:

1. **Re-title, don't retrain.** Position 20 years of production/systems work as "LLM systems in production." Build *one* shipped, measured RAG-or-agent feature as your proof artifact (Months 0–3).
2. **Go deep on the top-3 stack:** RAG (properly), evals/observability, agents+MCP. De-prioritize model training, framework tribalism, and anything AI commoditizes.
3. **Make agentic coding your daily driver** — it compounds your throughput *and* builds the marketable skill simultaneously.
4. **Specialize into a vertical you already know** by Months 9–18; aim to operate at the platform/AI-architecture altitude by Month 24 (highest value, lowest substitution).
5. **Build in public** — a measured write-up beats certifications as a hiring/credibility signal.

For the company (Innovation Ways-scale):

1. **Internal first:** adopt agentic coding + internal RAG + content/ops automation now; redesign *one* core workflow AI-native (the redesign is what creates ROI, per McKinsey).
 2. **Productize a vertical, not a wrapper:** pick one industry problem where your domain depth is the moat; target a 4–12 week MVP and harvest it from a real client POC.
 3. **Layer the services funnel** from R-00131/R-00135 (assessment → capped POC → build → retainer) with an EU-sovereignty/AI-Act premium; treat SME AI-Act readiness as a wedge offer.
 4. **Stay application/services-layer** — avoid capital-intensive, bubble-exposed infra bets; exploit falling token costs rather than owning compute.
-

IW Innovation Ways

CHAPTER 08

Limitations

CHAPTER 08

Limitations

- **Secondary-source skew.** Many statistics (SMB adoption %, agent-adoption data points, salary premiums) come from analyst/marketing aggregations rather than primary microdata; treat exact figures as **directional**. The most rigorous anchors here are McKinsey (EBIT/workflow), DORA 2025 (developer behavior), IEEE Spectrum/NACE (entry-level employment), and a16z (enterprise adoption order-of-magnitude).
 - **Fast-moving target.** Model capability, pricing, tooling, and the AI-Act timeline (the Digital Omnibus deferral is still *pending formal adoption*) can shift within the 12–24 month horizon; revisit quarterly.
 - **The bubble is a genuine tail risk.** A disorderly infra-layer correction would still ripple into application demand and hiring; the application/services-layer positioning recommended here is *more* resilient to it, not immune.
 - **Personalization gap.** This report does not know your specific domain background, financial runway, or risk appetite — the vertical-choice and services-vs-product decision should be made against those.
 - **Stanford "Enterprise AI Playbook" PDF** (51 deployments) was image-encoded and could not be parsed; its conclusions are corroborated here via a16z and McKinsey instead.
-

IW Innovation Ways

CHAPTER 09

Sources

CHAPTER 09

Sources

Cross-referenced prior research (IW AI Core DB): R-00129

(market landscape), R-00130 (SWE → AI pivot skills), R-00131 (IW GTM), R-00132 (solo showcase projects), R-00133 (agent eval/observability tooling), R-00134 (EU resume teardown), R-00135 (EU services pricing/packaging).

#	Title	Credibility	URL
1	a16z — Where Enterprises Are Actually Adopting AI	High (top-tier VC primary)	https://a16z.com/where-enterprises-are-actually-adopting-ai/
2	McKinsey — The State of AI	High (primary survey)	https://www.mckinsey.com/capabilities/quantumblack/our-insights/the-state-of-ai
3	DORA 2025 — State of AI-Assisted Software Development (Google Cloud)	High (rigorous primary)	https://cloud.google.com/blog/products/ai-machine-learning/announcing-the-2025-dora-report
4	IEEE Spectrum — AI Shifts Expectations for Entry-Level Jobs	High (incl. NACE 2026 data)	https://spectrum.ieee.org/ai-effect-entry-level-jobs
5	finalroundai — Software Engineering Job Market 2026	Medium (aggregator)	https://www.finalroundai.com/blog/software-engineering-job-market-2026
6	CNN Business — The demise of SWE jobs has been exaggerated	Medium-High (major outlet)	https://www.cnn.com/2026/04/08/tech/ai-software-developer-jobs
7	Digital Applied — AI Agent Adoption 2026 (120+ data points)	Medium (aggregator)	https://www.digitalapplied.com/blog/ai-agent-adoption-2026-enterprise-data-points
8	Man Group — The AI Bubble	High (institutional analysis)	https://www.man.com/insights/the-ai-bubble
9	tooldirectory — AI Capex Bubble 2026: where revenue actually is	Low-Medium (secondary)	https://tooldirectory.ai/blog/ai-capex-bubble-2026-where-the-revenue-actually-is
10	Futurum — AI Capex 2026: The \$690B Infrastructure Sprint	Medium (analyst)	https://futurumgroup.com/insights/ai-capex-2026-the-690b-infrastructure-sprint/
11	Security Boulevard — Why Senior Engineers Matter More in 2026	Medium (industry)	https://securityboulevard.com/2026/01/why-senior-software-engineers-will-matter-more-in-2026-in-an-ai-first-world/
12	Addy Osmani — The Next Two Years of Software Engineering	High (primary practitioner)	https://addyosmani.com/blog/next-two-years/
13	gnxtsystems — AI Skills 2025 obsolete, what 2026 demands	Low-Medium (secondary)	https://gnxtsystems.com/the-ai-skills-that-mattered-in-2025-are-already-obsolete-heres-what-2026-demands/
14	The New Stack — Cursor/Claude Code/Codex merging stack	Medium-High (industry press)	https://thenewstack.io/ai-coding-tool-stack/
15	Developers Digest — Best AI Coding Tools June 2026	Low-Medium (secondary)	https://www.developersdigest.tech/blog/best-ai-coding-tools-june-2026-post-fable5
16	MasterOfCode — AI ROI: Why Only 5% See Real Returns	Medium (cites MIT)	https://masterofcode.com/blog/ai-roi

17	Goldman Sachs — Small Businesses Embrace AI (survey)	High (primary survey)	https://www.goldmansachs.com/pressroom/press-releases/2026/small-businesses-embrace-ai-but-need-training-and-support-to-fully-harness-it
18	AdAI — Small Business AI Statistics 2026	Medium (aggregator)	https://adai.news/resources/statistics/small-business-ai-statistics-2026/
19	business.com — 2026 Small Business AI Outlook Report	Medium (primary-ish study)	https://www.business.com/articles/ai-usage-smb-workplace-study/
20	JPMorganChase Institute — Understanding AI use by small businesses	High (primary research)	https://www.jpmorganchase.com/institute/all-topics/business-growth-and-entrepreneurship/understanding-ai-use-by-small-businesses
21	Growtoria — Vertical Micro-SaaS: Why AI-Native Niche Wins 2026	Low-Medium (secondary)	https://growtoria.com/vertical-micro-saas-ai-businesses/
22	AI Magicx — Vertical AI Micro-SaaS Business Model 2026	Low-Medium (secondary)	https://www.aimagicx.com/blog/vertical-ai-micro-saas-business-model-2026
23	EU Commission — Regulatory Framework / AI Act	High (primary regulator)	https://digital-strategy.ec.europa.eu/en/policies/regulatory-framework-ai
24	artificialintelligenceact.eu — Small Businesses' Guide to the AI Act	Medium-High (specialist tracker)	https://artificialintelligenceact.eu/small-businesses-guide-to-the-ai-act/
25	Delbion — EU AI Act for SMEs: 7-Step Compliance Plan	Low-Medium (vendor)	https://www.delbion.com/en/insights/eu-ai-act-sme-compliance-guide/
26	futransolutions — Enterprise AI in 2026: co-creating ROI	Low-Medium (vendor)	https://futransolutions.com/blog/enterprise-ai-in-2026-how-business-and-it-co-create-real-roi/