

RESEARCH

Where Businesses Actually Get AI Value: Proven Use Cases and ROI by Function and Industry (2026)

Contents

01	Executive Summary	4
02	Background	7
03	Findings	10
04	Recommendations	17
05	Limitations	20
06	Sources	23

Research ID: R-00145

Date: 2026-06-23

Mode: market

Depth: deep

Primary Question: Across functions and industries, where is AI delivering measurable, repeatable ROI in 2026 — and what separates the deployments that pay off from the rest?

IW Innovation Ways

CHAPTER 01

Executive Summary

CHAPTER 01

Executive Summary

Enterprise AI investment tripled in a single year to [\\$37 billion in 2025](#), yet structural studies confirm that the vast majority of that spend produces no measurable financial return: MIT's NANDA initiative found [95% of enterprise AI deployments deliver no measurable P&L impact](#), and McKinsey found only [39% of organisations see any EBIT improvement — most of less than 5%](#). The value is real but radically concentrated: PwC's 2026 study places [74% of AI's economic gains inside just 20% of companies](#), and BCG finds high performers — defined as organisations with meaningful revenue or cash-flow increases — show three-year total shareholder returns roughly [four times higher](#) than laggards. The clearest separating factor is not model choice or budget size but **workflow redesign**: McKinsey identifies it as the single highest-impact factor for EBIT, and high performers are [nearly three times more likely to have fundamentally redesigned workflows](#) compared to the average enterprise.



IW Innovation Ways

CHAPTER 02

Background

CHAPTER 02

Background

This document builds on and extends two prior filed researches:

R-00129 established the 2026 AI market landscape: coding and customer support lead enterprise adoption, the ~95% pilot failure rate, Anthropic's 40% enterprise LLM market share, and model price deflation of ~90% since 2023.

R-00144 synthesised the a16z "barbell" model (coding leads by an order of magnitude, then support, then search), confirmed McKinsey's workflow-redesign differentiator, and covered high-level function and industry adoption patterns.

The gap R-00145 fills is concrete, function-by-function and industry-by-industry ROI data with named tools, actual numbers (payback periods, cost-per-unit savings, ARR benchmarks), the nuanced evidence on coding productivity (including the METR controlled-trial result that complicates the narrative), the failure pattern taxonomy in depth, and actionable guidance for a small EU software firm deciding where to apply or sell AI.

IW Innovation Ways

CHAPTER 03

Findings

CHAPTER 03

Findings

The headline macro numbers are now well-corroborated across multiple independent sources:

1. The Aggregate ROI Reality [HIGH confidence]

Metric	Figure	Source
Enterprise GenAI spend (2025)	\$37B (3x YoY)	Menlo Ventures
Pilots delivering no P&L impact	95%	MIT NANDA
Orgs reporting any EBIT impact	39%	McKinsey
Of those, EBIT >5% attributable to AI	~6% of all respondents	McKinsey
AI economic gains captured by top 20% of orgs	74%	PwC 2026
AI initiatives delivering expected ROI	25%	IBM CEO Study 2025
AI initiatives achieving enterprise-wide scale	16%	IBM CEO Study 2025
AI pilots converting to production (AI software)	47%	Menlo Ventures
Traditional software POC-to-production rate	~25%	Menlo Ventures
Agentic AI projects at cancellation risk by 2027	>40%	Gartner, June 2025

The paradox is structural: adoption is high (79% of organisations report regular GenAI use in at least one function per McKinsey), but the P&L signal remains narrow. The average payback period for a typical AI use case is [2–4 years](#), significantly longer than the 7–12 months CFOs expect from conventional technology investments, and [only 6% of deployments see returns within the first year](#).

What the ~5% Do Differently

McKinsey's high-performer profile (companies attributing >5% of EBIT to AI) shows four consistent patterns:

- Workflow redesign first:** [Nearly 3x more likely to have fundamentally redesigned workflows](#) before or alongside AI deployment, rather than bolting AI onto existing processes.
- Budget concentration:** [5x more likely to allocate over 20% of digital budgets to AI](#) — they treat AI as a transformation budget item, not a tooling line.
- Senior ownership:** [3x more likely to have strong senior leadership engagement](#) rather than delegating to a central AI lab.
- Agentic systems over chat interfaces:** Fewer than 10% of all organisations scale AI agents across any function; high performers are disproportionately in this group.

Deloitte corroborates the workflow-redesign finding: [only 16% of organisations have fully designed roles, processes, and operating models to integrate AI into work](#), and technology-first (not work-design-first) organisations are [1.6x more likely to report AI investments underperforming](#).

BCG summarises the gap bluntly: [only about 5% of organisations have managed to reap substantial financial gains from AI](#), and those companies show three-year total shareholder

returns roughly four times higher than AI laggards.

2. Value by Function [HIGH confidence for coding and support; MEDIUM for others]

2a. Software / Coding [MEDIUM confidence — controlled evidence is more nuanced than marketing claims]

Software development receives the largest slice of enterprise AI spend: [\\$4.0 billion in 2025 \(55% of departmental AI spend\)](#). GitHub Copilot holds ~42% of the AI coding assistant market; Cursor holds ~18%.

Controlled evidence is ambiguous:

- GitHub's own RCT found [tasks completed 55% faster, pull-request cycle time dropped from 9.6 to 2.4 days \(75% reduction\), and successful builds increased 84%](#). These are company-commissioned results on isolated tasks.
- METR's independent randomised controlled trial (July 2025, 16 experienced open-source developers, 246 real repository tasks) found the **opposite**: AI-assisted developers were [19% slower on average than the control group](#). Developers expected to be 24% faster and believed they had been 20% faster — a 39-point perception gap.
- Faros AI's telemetry study across [10,000 developers and 1,255 enterprise teams](#) found that high-AI-adoption teams merge 98% more pull requests, but PR review time increases 91%, and DORA metrics remain largely unchanged at the company level.

Synthesis: Individual throughput gains are real but context-dependent. Gains collapse at the company level without corresponding investment in code review infrastructure, testing pipelines, and release automation. The "10x engineer" narrative reflects the upper tail of performance on well-structured tasks, not an average enterprise outcome. Despite the nuance, 50% of enterprise developers now use AI daily per Menlo, and [token usage on AI coding tools is up 320x since 2023](#).

2b. Customer Support / Service [HIGH confidence]

Customer support is the second-largest function by AI investment and the clearest near-term ROI case. The economics are structurally clean: interactions are time-bounded, intents are constrained, and resolution is binary and verifiable.

- **Sierra:** Reached [\\$100M ARR in 21 months](#), then raised \$950M in May 2026 at a \$10B+ valuation. Pricing is outcome-based: enterprises pay per resolved interaction.
- **Intercom Fin:** Reports a [76% average resolution rate across 8,000+ customers](#), with over 40 million resolved conversations, at ~\$0.99 per resolved outcome.
- **Industry average:** Companies implementing AI in customer support reduce average cost per interaction by [68%, from \\$4.60 to \\$1.45](#), with an average reported return of \$3.50 per \$1 invested.

Why it works: output is verifiable, human escalation is structurally built in, and every resolved interaction improves future resolution rates.

2c. Sales and Marketing [MEDIUM confidence]

43% of sales reps actively use AI in 2025 (up from 24% in 2023 per HubSpot), and AI-assisted sellers report a 28% lift in prospect response rates. B2B companies using AI-powered lead generation report an average 73% increase in qualified leads within six months (vendor-adjacent), and HubSpot customers cite 129% more leads and 36% more deals closed in year one (conflates platform and AI effects). A structural headwind: AI search (Google AI Overviews, Perplexity) is collapsing traditional SEO channels, partially offsetting content-productivity gains.

2d. Internal Knowledge Search / Enterprise Search [MEDIUM confidence]

Glean is the category leader: from \$100M to \$200M ARR in nine months, \$7.2B valuation, with 40% wDAU/wMAU (twice the SaaS benchmark) and 100M+ agentic actions per year. ROI is primarily time-savings based; hard P&L attribution is difficult because the counterfactual (time wasted without search) is not tracked in ERP systems — which is why search ranks third, not first.

2e. Finance and Back-Office [MEDIUM confidence]

Finance shows the highest, most verifiable process-level ROI yet receives little of the AI budget. MIT NANDA flagged the misallocation: over half of GenAI budgets target sales/marketing, yet the biggest ROI exists in back-office automation.

Accounts Payable automation is the clearest case: manual invoice processing costs \$13–\$15; AI reduces it to under \$3 per invoice (76–80% reduction). A firm processing 2,000 invoices/month saves ~\$318,000/year; approval cycles compress from 19.5 to 3.2 days; most achieve 200–500% ROI in Year 1 with 2–6 month payback. It pays fast because unit economics are binary, data is structured, and the process is internally controlled.

2f. HR and People Operations [MEDIUM confidence]

AI-assisted recruiting delivers 20–40% cost-per-hire reduction and 30–50% time-to-fill improvement, saving recruiters 5–8 hours/week on scheduling. Vendor ROI claims (340% in 18 months) are uncontrolled. Only 12% of organisations have HR leading AI work-design, suggesting under-governance.

2g. Legal [HIGH confidence — market signal is strong]

Harvey crossed \$100M ARR in August 2025, an estimated \$300M by May 2026, used by 142,000+ lawyers across 1,500+ organisations, raising \$200M at an \$11B valuation. Power users save 36.9 hours/month (law firms) and 28.3 (in-house). Legal works because the work is

text-dense, high-repetition, measurable in time saved, and human-in-the-loop is already cultural — though [82% of legal AI leaders say assessing ROI is still a hurdle](#).

3. Value by Industry [MEDIUM-HIGH confidence]

3a. Financial Services [HIGH confidence]

[90% of financial institutions use AI for fraud detection](#); JPMorgan Chase has achieved [~\\$1.5B in AI cost savings](#) with false positives down 50% and 20% more fraud caught. AI fraud systems are projected to save global banks [over £9.6B annually by 2026](#); average AI ROI in financial services is [~180%, top performers 300%+](#).

3b. Healthcare [HIGH confidence — strongest evidence base]

Healthcare is the only industry with peer-reviewed evidence of AI ROI at scale, via ambient clinical documentation. Ambient scribes generated [\\$600M revenue in 2025](#). **Abridge** contracted [150+ health systems](#); UChicago Medicine research found clinicians spent [8.5% less total EHR time and 15%+ less time composing notes](#); a JAMA Network Open multicenter study found a [31% drop in burnout](#). It works because documentation time is directly measurable and EHR integration gives objective evidence independent of self-report.

3c. Retail [MEDIUM confidence]

Walmart's AI inventory optimisation improved [EBITDA margins by 2.3 points](#) (stockouts –25%, markdowns –18%); Target's personalisation raised [repeat purchase rates 22%](#); AI personalisation drives [2.5x engagement and a 31% conversion increase](#). Most results are from tier-1 retailers with massive proprietary data; mid-market faces a harder path.

3d. Manufacturing and Logistics [MEDIUM confidence]

UPS ORION saves [\\$400M annually](#); Lenovo's AI demand sensing cut [surplus inventory 20% and improved forecast accuracy 25%](#); industry average [23–31% cost reduction](#), but [only 6% see payback within one year](#) — most within 2–4 years, slowed by legacy integration.

3e. Technology Sector [HIGH confidence]

Tech leads all sectors; the [Ramp AI Index shows top AI spenders grew revenue 12% faster](#) than low-spending peers — the clearest revealed-preference ROI signal in the dataset.

4. The Success Pattern [HIGH confidence]

4a. Workflow redesign precedes or accompanies deployment. [McKinsey identifies it as the single highest-impact factor for EBIT among 25 attributes tested](#). Deloitte's 10-20-70 principle (70% people/process/culture) points the same way.

4b. Build vs. Buy: the market moved decisively to buying. [Share shifted from 53% purchased in 2024 to 76% in 2025](#); MIT found [purchasing succeeds ~67% vs. ~33% for internal builds](#).

4c. Vertical specificity beats horizontal generality. Harvey and Abridge outperform generic wrappers because they are trained on domain data and deeply integrated into workflows.

4d. Data architecture maturity as prerequisite. [68% of CEOs identify integrated data architecture as critical](#); discovering data gaps post-commitment incurs [60–120% of original project cost in remediation](#).

4e. Growth orientation over pure efficiency [MEDIUM]. [PwC finds capturing growth opportunities is the strongest factor in AI-driven financial performance](#), ahead of efficiency alone.

5. The Failure Pattern [HIGH confidence]

5a. Integration is the primary blocker. [47% of organisations cite technology integration limitations](#); AI built on clean pilot data fails on production's inconsistent schemas and missing APIs.

5b. Misallocated budgets toward high-visibility, low-ROI functions [MEDIUM]. [Over half of budgets target sales/marketing while the highest ROI is in back-office](#).

5c. Abandonment is accelerating. Enterprise AI abandonment jumped from [17% in 2024 to 42% in 2025](#), average sunk cost \$7.2M; [Gartner predicts >40% of agentic projects cancelled by 2027](#).

5d. The coding productivity paradox [MEDIUM]. [Individual output rises \(98% more PRs\) while DORA metrics stagnate because review bottlenecks increase 91%](#) — coding AI without lifecycle redesign redistributes work without eliminating it.

IW Innovation Ways

CHAPTER 04

Recommendations

CHAPTER 04

Recommendations

For a small EU software firm deciding where to apply or sell AI:

1. **Apply AI internally in customer support and back-office finance first.** Fastest, most verifiable payback (2–6 months for AP automation). Adopt a vertical-specific tool rather than building custom integrations.
 2. **For coding AI, invest in the whole lifecycle, not just the assistant.** Pair Copilot/Cursor/Claude Code with review tooling, test coverage, and deployment speed; measure DORA metrics before and after.
 3. **Sell into legal, healthcare, or financial services verticals with verifiable unit economics.** EU data sovereignty (GDPR-by-design) is a structural advantage US players struggle to match compliantly.
 4. **Lead with workflow redesign as a service, not just the tool.** Charging for change management improves unit economics and converts pilots to production (67% vs 33%).
 5. **Measure business outcomes from day one, not usage metrics.** [65% of high-ROI organisations prioritise use cases based on explicit outcome projections.](#) Define the counterfactual before deployment.
-

IW Innovation Ways

CHAPTER 05

Limitations

CHAPTER 05

Limitations

- **Self-report bias is pervasive** across McKinsey/IBM/Deloitte/BCG; MIT's 95%-no-impact finding is notable for running against it.
 - **Vendor sources are structurally inflated** (GitHub, Intercom, Abridge cite their own succeeding customers).
 - **The METR developer study is a small sample** (16 developers) — updates priors, doesn't close the book.
 - **EU-specific data is sparse**; EU adoption lags US ~12–18 months and faces AI Act compliance costs absent from US benchmarks.
 - **Payback estimates vary widely** with starting state and whether integration costs are included.
 - **Agentic AI ROI evidence is early**; most function-level evidence is for single-task AI.
-

IW Innovation Ways

CHAPTER 06

Sources

CHAPTER 06

Sources

Cross-referenced prior research (IW AI Core DB): R-00129 (market landscape), R-00144 (road-ahead capstone).

#	Title	Credibility	URL
1	McKinsey State of AI 2025	High	https://www.mckinsey.com/capabilities/quantumblack/our-insights/the-state-of-ai
2	McKinsey State of AI 2025 — CoLab analysis	High	https://www.colabsoftware.com/post/mckinseys-state-of-ai-2025-what-separates-high-performers-from-the-rest
3	MIT NANDA "GenAI Divide" — Fortune	High	https://fortune.com/2025/08/18/mit-report-95-percent-generative-ai-pilots-at-companies-failing-cfo/
4	Menlo Ventures 2025 State of GenAI in the Enterprise	High	https://menlovc.com/perspective/2025-the-state-of-generative-ai-in-the-enterprise/
5	Menlo Ventures 2025 State of AI in Healthcare	High	https://menlovc.com/perspective/2025-the-state-of-ai-in-healthcare/
6	PwC 2026 AI Performance Study	High	https://www.pwc.com/gx/en/news-room/press-releases/2026/pwc-2026-ai-performance-study.html
7	BCG — workflow redesign / high-performer TSR (Outsource Accelerator)	Medium	https://news.outsourceaccelerator.com/?p=77284
8	Deloitte: Work Design Essential to AI ROI	High	https://www.deloitte.com/us/en/about/press-room/work-design-essential-to-ai-roi.html
9	IBM CEO Study 2025	High	https://newsroom.ibm.com/2025-05-06-ibm-study-ceos-double-down-on-ai-while-navigating-enterprise-hurdles
10	a16z: Where Enterprises Are Actually Adopting AI	High	https://a16z.com/where-enterprises-are-actually-adopting-ai/
11	METR: Early-2025 AI on Experienced OSS Developer Productivity	High	https://metr.org/blog/2025-07-10-early-2025-ai-experienced-os-dev-study/
12	Faros AI: AI Productivity Paradox (10,000 developers)	High	https://www.faros.ai/ai-productivity-paradox
13	GitHub: Quantifying Copilot's Impact (RCT)	Medium	https://github.blog/news-insights/research/research-quantifying-github-copilots-impact-on-developer-productivity-and-happiness/
14	Gartner: >40% Agentic AI Projects Cancelled by 2027	High	https://www.gartner.com/en/newsroom/press-releases/2025-06-25-gartner-predicts-over-40-percent-of-agentic-ai-projects-will-be-canceled-by-end-of-2027
15	Glean: \$200M ARR press release	Medium	https://www.glean.com/press/glean-surpasses-200m-in-arr-for-enterprise-ai-doubling-revenue-in-nine-months
16	Glean \$7.2B valuation analysis	Medium	https://agentmarketcap.ai/blog/2026/04/07/glean-7b-enterprise-knowledge-search-agentic-category
17	Harvey: \$100M ARR (CNBC)	High	https://www.cnbc.com/2025/08/04/legal-ai-startup-harvey-revenue.html
18	Harvey: \$11B valuation (CNBC)	High	https://www.cnbc.com/2026/03/25/legal-ai-startup-harvey-raises-200-million-at-11-billion-valuation.html
19	Harvey: \$300M ARR estimate (Sacra)	Medium	https://sacra.com/c/harvey/

20	RSGL/Harvey adoption report	Medium	https://legaltechnology.com/2025/12/02/the-impact-of-legal-ai-a-deeper-dive-into-the-rsgi-harvey-adoption-report/
21	Harvey: Practical Framework for Legal AI ROI	Medium	https://www.harvey.ai/blog/practical-framework-for-legal-ai-roi
22	Sierra: \$100M ARR in 21 months (TechCrunch)	High	https://techcrunch.com/2025/11/21/bret-taylors-sierra-reaches-100m-arr-in-under-two-years/
23	Intercom Fin AI Guide	Medium	https://myaskai.com/blog/intercom-fin-ai-agent-complete-guide-2026
24	Abridge: Sharp HealthCare partnership	Medium	https://www.abridge.com/press-release/sharp-abridge
25	PMC: Ambient AI scribes review — JAMA burnout data	High	https://pmc.ncbi.nlm.nih.gov/articles/PMC12973079/
26	Advisory Board: Ambient AI and burnout	High	https://www.advisory.com/daily-briefing/2026/02/04/ambient-ai-oi-ec
27	PLANERGY: Accounts Payable 2025	Medium	https://planergy.com/blog/accounts-payable-in-2025/
28	Articsledge: AI in AP	Medium	https://www.articsledge.com/post/ai-accounts-payable-ap
29	Articsledge: AI fraud detection banking	Medium	https://www.articsledge.com/post/ai-fraud-detection-banking
30	Articsledge: Retail ML case studies	Medium	https://www.articsledge.com/post/machine-learning-retail-case-studies
31	MLVeda: AI supply chain guide	Medium	https://www.mlveda.com/blog/ai-native-supply-chain-complete-guide-to-intelligent-orchestration-roi
32	OpenSkyGroup: Supply chain AI statistics	Medium	https://openskygroup.com/supply-chain-ai-statistics/
33	UPS ORION savings (noloco)	Medium	https://noloco.io/blog/ai-in-logistics
34	HubSpot State of Sales 2025	Medium	https://blog.hubspot.com/sales/hubspot-sales-strategy-report
35	Cirrus Insight: AI in Sales 2025	Medium	https://www.cirrusinsight.com/blog/ai-in-sales
36	Humanly: AI recruiting cost reduction	Low-Medium	https://www.humanly.io/blog/reduce-recruiting-costs-ai-strategy
37	AptaHire: AI recruiting ROI breakdown	Low-Medium	https://aptahire.ai/ai-recruiting-cost-savings-breakdown/
38	Ramp AI Index	High	https://ramp.com/data/ai-index
39	FluidLabs: 42% AI project abandonment 2025	Medium	https://fluidlabs.com/resources/why-42-percent-enterprise-ai-abandoned-2025
40	CloudHew: Enterprise AI after pilot	Medium	https://cloudhew.com/blogs/why-enterprise-ai-projects-fail-after-the-pilot-the-hidden-cost-of-poor-cloud-data-and-integration-readiness/
41	AllAboutAI: Customer service AI statistics	Medium	https://www.allaboutai.com/resources/ai-statistics/customer-service/

42	Vellum AI: AI agent use cases guide	Medium	https://www.vellum.ai/blog/ai-agent-use-cases-guide-to-unlock-ai-roi
43	Digital Applied: AI ROI measurement framework 2026	Medium	https://www.digitalapplied.com/blog/ai-roi-measurement-framework-2026-cfo-payback-models
44	Caspian One: AI in Financial Services 2025	Medium	https://www.caspianone.com/ai-in-financial-services-report
45	Financial services AI ROI (thinking.inc)	Low-Medium	https://thinking.inc/en/industry-service/financial-services-ai-roi/
46	SuperAGI: Hyper-personalization in retail 2025	Low-Medium	https://superagi.com/hyper-personalization-in-retail-2025-how-ai-transforms-customer-journeys-with-real-time-predictive-analytics/