

RESEARCH

EU AI Act: Obligations Map for a Typical Company

Version v2 Status Draft Project IW AI Core Platform Generated 2026-06-24

Contents

01	Executive Summary	4
02	Introduction	7
03	Background & Context	10
04	Methodology	13
05	Findings	16
06	Analysis / Discussion	24
07	Conclusion	27
08	Recommendations	30
09	Limitations	33
10	Sources	36

Field	Value
ID	R-00154
Date	2026-06-24
Mode	deep
Depth	deep
Status	draft
Author	IW AI Core (/iw-research)
Series	EU AI Act Compliance Suite (1 of 4) — see R-00155 , R-00156 , R-00157

IW Innovation Ways

CHAPTER 01

Executive Summary

CHAPTER 01

Executive Summary

Bottom line: the EU AI Act (in force since 1 August 2024) imposes **graduated obligations driven by two variables: your role** (provider, deployer, importer, distributor, authorized representative) **and the risk tier** of each AI system (prohibited / high-risk / limited-transparency / minimal) [HIGH]. For most companies the live obligations are already here: the eight [Article 5 prohibitions](#) and the [Article 4 AI-literacy](#) duty (both since 2 Feb 2025), and GPAI provider rules (since 2 Aug 2025) [HIGH]. The **single most pressing upcoming deadline is 2 August 2026** — [Article 50 transparency](#) (chatbot disclosure, synthetic-content marking, deepfake labelling) — which the **Digital Omnibus did *not* defer** [HIGH]. The Digital Omnibus (EP-adopted 16 June 2026, 423–57; Council adoption expected 29 June 2026) **did** push standalone high-risk (Annex III) to **2 Dec 2027** and product-embedded high-risk (Annex I) to **2 Aug 2028** [HIGH], but is **not yet legally in force**, so the original timeline remains the binding baseline until Official-Journal publication [MEDIUM].



IW Innovation Ways

CHAPTER 02

Introduction

CHAPTER 02

Introduction

This is document 1 of a four-part EU AI Act suite. A typical company — most often a deployer buying AI tools, sometimes a provider building them — needs a clear, role-aware answer to one practical question: *what actually applies to us, and by when?* This document supplies that obligations map as framing for the rest of the suite.

Why this research

Objectives

The primary question driving this research: **what obligations does the EU AI Act (Regulation (EU) 2024/1689) actually impose on a *typical company*, depending on its role in the AI supply chain and the risk tier of its AI systems — and by when?** The objective is to produce a definitive, dated, role-and-tier obligations map a typical company can use to determine its exposure.

Scope

This document is the **obligations map**: which obligations attach to which role and risk tier, and on what date each becomes binding. The companion documents are siblings: the controls to implement and how to test them ([R-00155](#)), the step-by-step compliance roadmap ([R-00156](#)), and the business opportunity of selling compliance services ([R-00157](#)).

IW Innovation Ways

CHAPTER 03

Background & Context

CHAPTER 03

Background & Context

The EU AI Act is **risk-proportionate**: the higher the potential harm to health, safety, or fundamental rights, the heavier the burden — ranging from outright bans, through extensive provider/deployer duties for high-risk systems, down to transparency-only or no obligations for the rest. The weight of obligation for any given system is therefore a function of *two* variables only — the risk tier the system lands in, and the legal role the company plays toward it.

The Act applies **extraterritorially**: it binds any provider placing a system on the EU market and any deployer using one in the EU, regardless of where the company is established. This is why even non-EU companies serving EU users fall in scope.

This is document 1 of a four-part suite. It maps *what applies to whom and when*; the companion documents cover the controls to implement and how to test them ([R-00155](#)), the step-by-step roadmap ([R-00156](#)), and the business opportunity of selling compliance services ([R-00157](#)).

IW Innovation Ways

CHAPTER 04

Methodology

CHAPTER 04

Methodology

This research is **primary-source-anchored**. Obligation content, dates, and thresholds were drawn first from the primary legal text (Regulation (EU) 2024/1689 on EUR-Lex), the annotated AI Act Explorer, and official Commission / Council / AI Act Service Desk material; law-firm analyses were used to corroborate dating and the Digital Omnibus legislative trajectory, not to establish law on their own.

Confidence is declared per finding on a three-point scale:

- **HIGH** — supported by primary law or official sources, multiply corroborated.
 - **MEDIUM** — supported but contingent (e.g., not-yet-in-force legislation, outstanding implementing acts).
 - **LOW** — single-source or unverified in retrieved material.
-

IW Innovation Ways

CHAPTER 05

Findings

CHAPTER 05

Findings

Figure 1. Risk-tier classification path for any AI system. Most business tools land in Minimal or Limited risk (source: [Art. 5](#), [Art. 6](#), [Annex III](#)).

F1 — The four risk tiers determine the weight of obligation [HIGH]

Why: every obligation downstream depends first on which tier a system lands in — this is the entry point of the whole regime.

Prohibited (Art. 5): eight non-remediable practices — subliminal manipulation, exploiting vulnerabilities, social scoring, criminal-risk profiling, untargeted facial scraping, workplace/education emotion inference, sensitive biometric categorization, and real-time remote biometric ID in public spaces. **High-risk:** Annex I product-safety components (medical devices, machinery, vehicles...) and eight Annex III standalone categories (biometrics, critical infrastructure, education, employment, essential services, law enforcement, migration, justice). **Limited:** chatbots and generative systems — transparency only. **Minimal:** spam filters, recommenders, analytics — no obligations [HIGH] ([high-level summary](#), [Commission overview](#)).

F2 — Your role decides which duties you bear; deployers get no SME exemption from Art. 26 [HIGH]

The Act defines five roles in [Article 3](#). The two most companies occupy:

- **Provider** — develops/markets an AI system under its own name. Bears the heaviest load for high-risk: QMS (Art. 17), technical docs (Art. 11), risk management (Art. 9), data governance (Art. 10), logging (Art. 12), human oversight (Art. 14), accuracy/robustness/security (Art. 15), conformity assessment (Arts. 43–46), CE marking (Art. 48), EU-database registration (Art. 49) ([Art. 16](#)).
- **Deployer** — uses an AI system under its own authority (the typical SME buying a tool). For high-risk: follow instructions, assign human oversight, monitor, retain logs ≥6 months, notify workers, inform affected persons, and (public bodies/public-service providers) run a Fundamental Rights Impact Assessment ([Art. 26](#), [Art. 27]). **Critically, Article 26 has no SME carve-out** [HIGH].

Warning

Role-shifting (Art. 25) is silent and automatic. A deployer becomes a provider — with full provider obligations — the moment it puts its own brand on a high-risk system, substantially modifies one, or repurposes any system so it becomes high-risk. No registration step; the legal status simply changes (Art. 25).

F3 — Article 5 prohibitions are live (since 2 Feb 2025); two new ones arrive 2 Dec 2026 [HIGH]

The eight prohibitions became enforceable 2 Feb 2025, with [Commission guidelines](#) published 4 Feb 2025; penalties activated 2 Aug 2025 ([Mayer Brown](#)). The Digital Omnibus adds **two new prohibitions effective 2 Dec 2026**: AI generating non-consensual intimate imagery (NCII) and child sexual abuse material (CSAM) ([Gibson Dunn](#)).

F4 — Article 50 transparency is the urgent near-term deadline (2 Aug 2026) [HIGH]

Four cross-tier disclosure duties, applying regardless of risk tier ([Art. 50](#), [AI Act Service Desk](#)):

Provision	Duty	Burden
50(1)	Tell users they're interacting with AI, at first interaction	Provider
50(2)	Mark synthetic audio/image/video/text machine-readable	Provider
50(3)	Disclose emotion-recognition / biometric categorization to exposed persons	Deployer
50(4)	Disclose deepfakes & AI-generated public-interest text	Deployer

50(2) marking applies immediately to systems placed on market from 2 Aug 2026; **pre-existing systems get a grace period to 2 Dec 2026**. The Commission issued a draft [Code of Practice on AI-Generated Content](#) on 10 June 2026 ([Greenberg Traurig](#)). **Not deferred by the Omnibus** [HIGH].

F5 — Article 4 AI literacy applies to everyone using AI; supervision starts 2 Aug 2026 [HIGH]

Every provider and deployer must ensure "a sufficient level of AI literacy" of staff — proportionate, role-differentiated, no certificate required, but "relying on instructions for use" is insufficient ([Art. 4, Commission AI-literacy Q&A](#)). In force since 2 Feb 2025; authority supervision begins 2 Aug 2026 ([Latham & Watkins](#)).

F6 — GPAI model providers have been regulated since 2 Aug 2025 [HIGH]

GPAI providers must keep technical documentation (Annex XI), give downstream integrators capability/limitation info (Annex XII), publish an EU-copyright policy honouring rights reservations, and publish a training-data summary ([Art. 53, Commission GPAI FAQ](#)). Models with **systemic risk** (legal presumption at **10²⁵ FLOPs** training compute) add evaluation, adversarial testing, risk mitigation, incident reporting, and cybersecurity ([Art. 55](#)). The final [GPAI Code of Practice](#) (10 July 2025) grants a presumption of conformity. Open-source models are exempt from documentation duties but **not** from copyright/training-summary duties, and the exemption never applies to systemic-risk models.

F7 — High-risk obligations are extensive but mostly deferred to 2027–2028 [HIGH]

Providers of high-risk systems must satisfy Articles 9–17 (risk management, data governance, technical documentation, logging, deployer transparency, human oversight, accuracy/robustness/security, QMS); deployers satisfy Article 26 ([A&O Shearman](#)). Most Annex III systems self-assess (Annex VI); **biometric** systems and certain categories without harmonized standards require a third-party **notified body** (Annex VII). Implementation detail lives in [R-00155](#).

F8 — The full enforcement timeline (with Digital Omnibus revisions) [HIGH]

Why: the staggered dates are the planning backbone — what's overdue, what's imminent, and what you have runway for.




Figure 2. Enforcement schedule as of June 2026. Article 50 (Aug 2026) is the next hard deadline; high-risk obligations were pushed to 2027–2028 by the Digital Omnibus (sources: [AI Act Service Desk timeline](#), [implementation timeline](#), [Sidley](#)).

F9 — Penalties: up to €35M / 7% turnover — but SMEs pay the *lower* figure [HIGH]

Three tiers under [Article 99](#), on **worldwide** turnover:

Tier	Cap	Covers
1	€35M or 7%	Art. 5 prohibitions; GPAL systemic-risk
2	€15M or 3%	Provider (Art. 16), deployer (Art. 26), Art. 50, notified-body duties
3	€7.5M or 1%	Incorrect/misleading information to authorities

The **SME inversion (Art. 99(6))**: large firms pay *whichever is higher*; SMEs/startups pay *whichever is lower*. A €2M-turnover startup breaching Art. 50 faces a max of €60K (3%), not €15M; a €10M startup breaching a prohibition faces €700K (7%), not €35M [HIGH] ([Legalithm](#)). (Your seed example used €140K for a €2M firm at Tier 1, which is 7% of €2M; the math holds for a Tier-1 breach.)

F10 — Digital Omnibus: confirmed, but not yet legally in force [HIGH political / MEDIUM legal]

Why: your runway for high-risk systems hinges entirely on whether this is law yet — and as of today it is not.

Your seed facts are **confirmed and sharpened**: provisional agreement **7 May 2026** (after a failed 28 Apr trilogue), Council confirmation 13 May, **EP plenary adoption 16 June 2026 (423–57, 174 abstentions)**, Council formal adoption expected **29 June 2026**, OJ publication expected July 2026 ([Council press release](#), [iubenda](#), [White & Case](#)). Confirmed changes on publication: Annex III → 2 Dec 2027; Annex I → 2 Aug 2028; sandboxes → 2 Aug 2027; watermarking grace → 2 Dec 2026; new NCII/CSAM bans; a GDPR tweak allowing sensitive

data for bias mitigation; broadened SME/“small mid-cap” relief ([Dastra](#), [Inside Privacy](#)). **Until OJ publication, the original 2 Aug 2026 Annex III deadline is the binding baseline** [MEDIUM].

IW Innovation Ways

CHAPTER 06

Analysis / Discussion

CHAPTER 06

Analysis / Discussion

The findings resolve into a clean live-versus-deferred split that should drive every company's prioritization. The obligations already biting — Art. 5 prohibitions (F3), Art. 4 AI literacy (F5), and GPAI provider rules (F6) — are non-negotiable today and apply irrespective of risk tier or company size; Art. 26 in particular gives deployers no SME relief (F2). The next hard wall is **Article 50 transparency on 2 Aug 2026** (F4), and its significance is sharpened by F10: the Digital Omnibus deliberately did *not* touch Art. 50, so the disclosure/marking duties land on schedule even as the heavier high-risk regime slides. That contrast — Art. 50-now versus high-risk-2027/2028 — is the single most important planning takeaway, because it inverts the intuition that "high-risk is the urgent part."

The second cross-cutting theme is legal contingency. The deferrals to 2 Dec 2027 (Annex III) and 2 Aug 2028 (Annex I), the new NCII/CSAM bans, and the broadened SME relief are all products of an Omnibus that is **adopted politically but not yet in force** (F10), meaning the original 2 Aug 2026 high-risk baseline remains technically binding until OJ publication. Layered on top, the role-shift trap (F2) and the SME penalty inversion (F9) mean a company's true exposure is a moving function of role, turnover, and the system it touches — not a fixed checklist. A company that fine-tunes or rebrands a third-party high-risk system can silently acquire the full provider load, while a small firm's downside on most breaches is capped at the *lower* of the two penalty figures.

IW Innovation Ways

CHAPTER 07

Conclusion

CHAPTER 07

Conclusion

The EU AI Act imposes on a typical company a graduated obligation set determined jointly by its role (most often deployer, sometimes provider) and the risk tier of each AI system it builds or uses — with the prohibitions, AI-literacy duty, and GPAI rules already live, Article 50 transparency due 2 August 2026, and the bulk of high-risk obligations deferred to December 2027 and August 2028. For the typical company, the practical reading is that the urgent work is not high-risk compliance but the here-and-now: confirm no system trips an Art. 5 prohibition, stand up an AI-literacy program, and prepare Article 50 disclosures for the 2 August 2026 deadline that the Digital Omnibus pointedly left in place. High-risk runway exists, but it is a deferral — not an exemption — and it is not yet even legally enacted, so any company touching an Annex III use case should treat the gap analysis as starting now. Role and size change the picture materially: a rebrand or substantial modification can convert a deployer into a fully-burdened provider, while an SME's penalty exposure is capped at the lower of the two statutory figures.



CHAPTER 08

Recommendations

CHAPTER 08

Recommendations

1. **Now (overdue):** Audit every AI system against the eight Art. 5 prohibitions — especially HR emotion inference and any biometric tooling.
 2. **Now (in force):** Stand up a proportionate, documented AI-literacy program (Art. 4).
 3. **By 2 Aug 2026 (urgent):** Implement Article 50 disclosures and confirm every AI vendor builds in disclosure/marketing. Review public-facing AI-content workflows for 50(4).
 4. **By 2 Dec 2026:** Close the watermarking grace period for pre-existing generative systems; eliminate any NCII/CSAM-capable features.
 5. **Through 2 Dec 2027:** If you touch any Annex III use case (recruitment, credit, essential services, education), begin gap analysis now — the obligation set is large even with the deferral.
 6. **GPAI providers:** Verify copyright policy, training-data summary, and documentation; check the 10^{25} FLOPs systemic-risk threshold.
 7. **Watch the OJ:** Don't bank the 2 Dec 2027 extension until the Omnibus is published; the 2 Aug 2026 Annex III date is technically still law today.
 8. **Run a role audit:** Anyone re-branding/fine-tuning/modifying third-party AI should check the Art. 25 provider-shift triggers.
-

IW Innovation Ways

CHAPTER 09

Limitations

CHAPTER 09

Limitations

- **Digital Omnibus not yet in force (as of 24 June 2026)** — original timeline remains the legal baseline until OJ publication; the deferred dates are high-probability, not enacted.
 - **Enforcement is unproven** — national authorities were designated from Aug 2025 but capacity varies across 27 states; no major penalty decisions reported yet.
 - **Art. 50 implementing acts / final Code of Practice** on synthetic-content marking remain outstanding (draft only, 10 June 2026).
 - **Art. 6(3) high-risk classification guidelines** (due 2 Feb 2026) — publication/content not confirmed in retrieved sources.
 - **SME vs "small mid-cap" boundary** under the Omnibus needs verification against the final OJ text.
-

IW Innovation Ways

CHAPTER 10

Sources

CHAPTER 10

Sources

#	Title	Credibility	URL
1	Regulation (EU) 2024/1689 — EUR-Lex (OJ)	High (primary law)	https://eur-lex.europa.eu/eli/reg/2024/1689/oj/eng
2	AI Act Explorer — Art. 3 Definitions	High (primary, annotated)	https://artificialintelligenceact.eu/article/3/
3	AI Act Explorer — Art. 5 Prohibited Practices	High	https://artificialintelligenceact.eu/article/5/
4	AI Act Explorer — Art. 6 Classification	High	https://artificialintelligenceact.eu/article/6/
5	AI Act Explorer — Art. 16 Provider Obligations	High	https://artificialintelligenceact.eu/article/16/
6	AI Act Explorer — Art. 25 Value Chain	High	https://artificialintelligenceact.eu/article/25/
7	AI Act Explorer — Art. 26 Deployer Obligations	High	https://artificialintelligenceact.eu/article/26/
8	AI Act Explorer — Art. 50 Transparency	High	https://artificialintelligenceact.eu/article/50/
9	AI Act Explorer — Art. 53 GPAI Obligations	High	https://artificialintelligenceact.eu/article/53/
10	AI Act Explorer — Art. 55 Systemic Risk	High	https://artificialintelligenceact.eu/article/55/
11	AI Act Explorer — Art. 99 Penalties	High	https://artificialintelligenceact.eu/article/99/
12	AI Act Explorer — Annex III	High	https://artificialintelligenceact.eu/annex/3/
13	AI Act Explorer — Implementation Timeline	High	https://artificialintelligenceact.eu/implementation-timeline/
14	AI Act Explorer — High-Level Summary	High	https://artificialintelligenceact.eu/high-level-summary/
15	AI Act Service Desk — Art. 50 (Commission)	High (official)	https://ai-act-service-desk.ec.europa.eu/en/ai-act/article-50
16	AI Act Service Desk — Art. 99 (Commission)	High (official)	https://ai-act-service-desk.ec.europa.eu/en/ai-act/article-99
17	AI Act Service Desk — Implementation Timeline	High (official)	https://ai-act-service-desk.ec.europa.eu/en/ai-act/timeline/timeline-implementation-eu-ai-act
18	Commission — AI Literacy Q&A	High (official)	https://digital-strategy.ec.europa.eu/en/faqs/ai-literacy-questions-answers
19	Commission — Guidelines on Prohibited AI Practices	High (official)	https://digital-strategy.ec.europa.eu/en/library/commission-publishes-guidelines-prohibited-artificial-intelligence-ai-practices-defined-ai-act
20	Commission — Code of Practice on AI-Generated Content	High (official)	https://digital-strategy.ec.europa.eu/en/policies/code-practice-ai-generated-content
21	Commission — GPAI FAQ	High (official)	https://digital-strategy.ec.europa.eu/en/faqs/general-purpose-ai-models-ai-act-questions-answers

22	Commission — AI Act policy page	High (official)	https://digital-strategy.ec.europa.eu/en/policies/regulatory-framework-ai
23	Council press release — 7 May 2026 Omnibus agreement	High (primary)	https://www.consilium.europa.eu/en/press/press-releases/2026/05/07/artificial-intelligence-council-and-parliament-agree-to-simplify-and-streamline-rules/
24	iubenda — Parliament adoption June 2026	Medium-High	https://www.iubenda.com/en/blog/ai-omnibus-parliament-adoption-june-2026/
25	Sidley — Digital Omnibus analysis (22 Jun 2026)	Medium-High (law firm)	https://datamatters.sidley.com/2026/06/22/eu-lawmakers-reach-provisional-agreement-to-delay-key-eu-ai-act-obligations/
26	Gibson Dunn — Omnibus postponed deadlines	Medium-High (law firm)	https://www.gibsondunn.com/eu-ai-act-omnibus-agreement-postponed-high-risk-deadlines-and-other-key-changes/
27	Dastra — Omnibus Parliament vote details	Medium	https://www.dastra.eu/en/blog/digital-omnibus-on-ai-parliament-votes-deadlines-redrawn/60108
28	Inside Privacy — Digital Omnibus May 2026	Medium-High (law firm)	https://www.insideprivacy.com/artificial-intelligence/eu-ai-act-update-timeline-relief-targeted-simplification-and-new-prohibitions/
29	White & Case — Digital Omnibus deal	Medium-High (law firm)	https://www.whitecase.com/insight-alert/eu-agrees-digital-omnibus-deal-simplify-ai-rules
30	Latham & Watkins — GPAI obligations in force	Medium-High (law firm)	https://www.lw.com/en/insights/eu-ai-act-gpai-model-obligations-in-force-and-final-gpai-code-of-practice-in-place
31	A&O Shearman — high-risk obligations	Medium-High (law firm)	https://www.aoshearman.com/en/insights/ao-shearman-on-tech/zooming-in-on-ai-10-eu-ai-act-what-are-the-obligations-for-high-risk-ai-systems
32	Mayer Brown — Art. 5 & Art. 4 in effect	Medium-High (law firm)	https://www.mayerbrown.com/en/insights/publications/2025/01/eu-ai-act-ban-on-certain-ai-practices-and-requirements-for-ai-literacy-come-into-effect
33	Greenberg Traurig — Art. 50 (June 2026)	Medium-High (law firm)	https://www.gtlaw.com/en/insights/2026/6/deepfakes-chatbots-ai-generated-text-european-commission-details-transparency-obligations-under-the-ai-act
34	Latham & Watkins — Art. 4 training & prohibitions	Medium-High (law firm)	https://www.lw.com/en/insights/upcoming-eu-ai-act-obligations-mandatory-training-and-prohibited-practices
35	Legalithm — SME/startup compliance guide	Medium	https://www.legalithm.com/en/blog/eu-ai-act-compliance-guide-startups-smes